

(ubuntu 18.04 cuda 10.2 pytorch onnx→tensorrt) 使用tensorrt加载并运行onnx模型

2021.4

0. TensorRT介绍

1. 安装tensorrt、 pycuda

1.1 安装tensorrt

1.2 安装pycuda

2. 将pytorch模型转化为onnx模型

3. tensorrt调用onnx模型进行预测

4. 注意事项


- 下述代码链接：

https://github.com/kiyoxi2020/tensorrt_example

0. TensorRT介绍

<https://docs.nvidia.com/deeplearning/tensorrt/developer-guide/index.html>

- TensorRT的核心是一个C ++库，可促进对NVIDIA图形处理单元（GPU）进行高性能推断。它旨在与TensorFlow，Caffe，PyTorch，MXNet等深度学习框架互补地工作。它专门致力于在GPU上快速有效地运行已经受过训练的网络，以生成结果（包括检测，回归或推断任务等）。
- TensorRT根据指定的精度（FP32，FP16或INT8）组合各层，优化内核选择，通过归一化、转换为更优化的矩阵数学运算等方式，改善延迟，吞吐量和效率。
- 主要流程：



1. 安装tensorrt、 pycuda

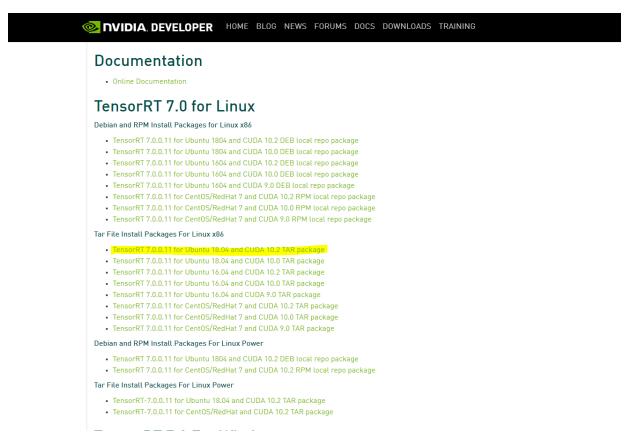
1.1 安装tensorrt

- 环境：ubuntu18.04, 2080Ti, 安装cuda10.2+cuda 7.6.5, 且在~/.bashrc中添加cuda、cudnn所在位置：

```
export CUDA_HOME=/usr/local/cuda
export CUDNN_INCLUDE_PATH=/usr/local/cuda/include
export CUDNN_LIB_DIR=/usr/local/cuda/lib64/
export CUDNN_PATH=/usr/local/cuda/lib64/libcudnn.so
export CUDNN_LIBRARY=/usr/local/cuda/lib64
export PATH=$PATH:$CUDA_HOME/bin
export LD_LIBRARY_PATH=/usr/local/cuda-10.2/lib64:$LD_LIBRARY_PATH
```

- tensorrt下载链接：<https://developer.nvidia.com/zh-cn/tensorrt>

需要选择tensorrt 7.0.0.11（适用于Ubuntu18.04、cuda10.2、cudnn7.6，其他更高的tensorrt版本要求cudnn8，因此此处选择tensorrt 7.0.0.11）



- 下载之后解压：

```
tar xzvf TensorRT-7.0.0.11.Ubuntu-18.04.x86_64-gnu.cuda-10.2.cudnn7.6.tar.gz
```

- 解压后得到“TensorRT-7.0.0.11”文件，包含以下目录：

```
bin -> targets/x86_64-linux-gnu/bin
data
doc
graphsurgeon
include
lib -> targets/x86_64-linux-gnu/lib
python
samples
targets
TensorRT-Release-Notes.pdf -> doc/pdf/TensorRT-Release-Notes.pdf
uff
10 directories, 1 file
```

- 安装指导：打开[doc/pdf/TensorRT-Installation-Guide.pdf](#)

由于下载的是tar安装文件，因此根据4.3 tar file installation安装即可

主要安装步骤如下:

```
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:~/TensorRT-7.0.0.11/lib #TensorRT文件所在目录
cd TensorRT-7.0.0.11/python
pip3 install tensorrt-* -cp3x-none-linux_x86_64.whl # 根据python版本选择whl文件

cd TensorRT-7.0.0.11/graphsurgeon
pip3 install graphsurgeon-0.4.1-py2.py3-none-any.whl
```

- 验证安装：进入samples目录，该目录下有很多c++、python的tensorrt使用示例，可以详细学习一下
 - 进入samples/sampleMNIST目录

```
make clean    # 可能需要sudo
make

cd ../../bin
./sample_mnist
```

- 还需要准备数据：进入data/minist目录下，运行

```
python download_pgms.py
```

得到0.pgm~9.pgm，注意download_pgms.py里面的mnist下载链接可能无效，此时可以找一个有效的链接进行替换

- 得到结果

[illegible]

1.2 安装pycuda

```
pip install pycuda
```

2. 将pytorch模型转化为onnx模型

- tensorrt无法直接读取pytorch模型，因此需要先转化为onnx
- 以alexnet为例，主要代码：

```
import torch
import torchvision

def get_model():
    model = torchvision.models.alexnet(pretrained=True).cuda()
    resnet_model = model.eval()
    return model

def get_onnx(model, onnx_save_path, example_tensor):
    example_tensor = example_tensor.cuda()
    _ = torch.onnx.export(model,
                          example_tensor,
                          onnx_save_path,
                          verbose=True,
                          training=False,
                          do_constant_folding=False,
                          input_names=['input'],
                          output_names=['output'])

if __name__ == '__main__':
    model = get_model()
    onnx_save_path = "alexnet.onnx"
    example_tensor = torch.randn(1, 3, 224, 224, device='cuda')

    get_onnx(model, onnx_save_path, example_tensor)
```

会有onnx模型的输出：

```

graph(%input : Float(1:150528, 3:50176, 224:224, 224:1),
  %features.0.weight : Float(64:363, 3:121, 11:11, 11:1),
  %features.0.bias : Float(64:1),
  %features.3.weight : Float(102:1600, 64:25, 5:5, 5:1),
  %features.3.bias : Float(102:1),
  %features.6.weight : Float(384:1728, 192:9, 3:3, 3:1),
  %features.6.bias : Float(384:1),
  %features.8.weight : Float(256:3456, 384:9, 3:3, 3:1),
  %features.8.bias : Float(256:1),
  %features.10.weight : Float(256:2304, 256:9, 3:3, 3:1),
  %features.10.bias : Float(256:1),
  %classifier.1.weight : Float(4096:9216, 9216:1),
  %classifier.1.bias : Float(4096:1),
  %classifier.4.weight : Float(4096:4096, 4096:1),
  %classifier.4.bias : Float(4096:1),
  %classifier.6.weight : Float(1000:4096, 4096:1),
  %classifier.6.bias : Float(1000:1))
%17 : Float(1:193600, 64:3025, 55:55, 55:1) = onnx::Conv[dilations=[1, 1], group=1, kernel_shape=[11, 11], pads=[2, 2, 2, 2], strides=[4, 4]](%input, %features.0.weight, %features.0.bias) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/modules/conv.py:416:0
%18 : Float(1:193600, 64:3025, 55:55, 55:1) = onnx::Relu(%17) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1117:0
%19 : Float(1:46656, 64:729, 27:27, 27:1) = onnx::MaxPool[kernel_shape=[3, 3], pads=[0, 0, 0, 0], strides=[2, 2]](%18) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:576:0
%20 : Float(1:139968, 192:729, 27:27, 27:1) = onnx::Conv[dilations=[1, 1], group=1, kernel_shape=[5, 5], pads=[2, 2, 2, 2], strides=[1, 1]](%19, %features.3.weight, %features.3.bias) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/modules/conv.py:416:0
%21 : Float(1:139968, 192:729, 27:27, 27:1) = onnx::Relu(%20) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1117:0
%22 : Float(1:32448, 192:169, 13:13, 13:1) = onnx::MaxPool[kernel_shape=[3, 3], pads=[0, 0, 0, 0], strides=[2, 2]](%21) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:576:0
%23 : Float(1:64896, 384:169, 13:13, 13:1) = onnx::Conv[dilations=[1, 1], group=1, kernel_shape=[3, 3], pads=[1, 1, 1, 1], strides=[1, 1]](%22, %features.6.weight, %features.6.bias) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/modules/conv.py:416:0
%24 : Float(1:64896, 384:169, 13:13, 13:1) = onnx::Relu(%23) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1117:0
%25 : Float(1:43264, 256:169, 13:13, 13:1) = onnx::Conv[dilations=[1, 1], group=1, kernel_shape=[3, 3], pads=[1, 1, 1, 1], strides=[1, 1]](%24, %features.8.weight, %features.8.bias) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/modules/conv.py:416:0
%26 : Float(1:43264, 256:169, 13:13, 13:1) = onnx::Relu(%25) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1117:0
%27 : Float(1:43264, 256:169, 13:13, 13:1) = onnx::Conv[dilations=[1, 1], group=1, kernel_shape=[3, 3], pads=[1, 1, 1, 1], strides=[1, 1]](%26, %features.10.weight, %features.10.bias) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/modules/conv.py:416:0
%28 : Float(1:43264, 256:169, 13:13, 13:1) = onnx::Relu(%27) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1117:0
%29 : Float(1:9216, 256:36, 6:6, 6:1) = onnx::MaxPool[kernel_shape=[3, 3], pads=[0, 0, 0, 0], strides=[2, 2]](%28) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:576:0
%30 : Float(1:9216, 256:36, 6:6, 6:1) = onnx::AveragePool[kernel_shape=[1, 1], strides=[1, 1]](%29) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:926:0
%31 : Float(1:9216, 9216:1) = onnx::Flatten[axis=1](%30) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:973:0
%32 : Float(1:4096, 4096:1) = onnx::Gemm[alpha=1., beta=1., transB=1](%31, %classifier.1.weight, %classifier.1.bias) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1674:0
%33 : Float(1:4096, 4096:1) = onnx::Relu(%32) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1117:0
%34 : Float(1:4096, 4096:1) = onnx::Gemm[alpha=1., beta=1., transB=1](%33, %classifier.4.weight, %classifier.4.bias) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1674:0
%35 : Float(1:4096, 4096:1) = onnx::Relu(%34) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1117:0
%output : Float(1:1000, 1000:1) = onnx::Gemm[alpha=1., beta=1., transB=1](%35, %classifier.6.weight, %classifier.6.bias) # /home/wangyuqing/anaconda3/envs/torch-gpu/lib/python3.6/site-packages/torch/nn/functional.py:1674:0
return (%output)

```

3. tensorrt调用onnx模型进行预测

- 参考1: <https://zhuanlan.zhihu.com/p/88318324>
- 参考2: samples/python/yolov3_onnx
- 参考3: doc/pdf/TensorRT-Developer-Guide.pdf (Using The Python API章节)
- 主要代码:

```

import pycuda.autoint
import numpy as np
import pycuda.driver as cuda
import tensorrt as trt
import torch
import os
import time
import torchvision

max_batch_size = 1
onnx_model_path = 'alexnet.onnx'
trt_engine_path = 'alexnet.trt'

TRT_LOGGER = trt.Logger()

class HostDeviceMem(object):
    def __init__(self, host_mem, device_mem):
        self.host = host_mem
        self.device = device_mem

    def __str__(self):
        return "Host: \n" + str(self.host) + '\nDevice:\n' + str(self.device)

    def __repr__(self):
        return self.__str__()

```

```

def allocate_buffers(engine):
    inputs = []
    outputs = []
    bindings = []
    stream = cuda.Stream()
    for binding in engine:
        size = trt.volume(engine.get_binding_shape(binding)) * engine.max_batch_size
        dtype = trt.nptype(engine.get_binding_dtype(binding))
        host_mem = cuda.pagelocked_empty(size, dtype)
        device_mem = cuda.mem_alloc(host_mem.nbytes)
        bindings.append(int(device_mem))
        if engine.binding_is_input(binding):
            inputs.append(HostDeviceMem(host_mem, device_mem))
        else:
            outputs.append(HostDeviceMem(host_mem, device_mem))

    return inputs, outputs, bindings, stream

def get_engine(max_batch_size=1, onnx_file_path="", engine_file_path="", save_engine=False):

    def build_engine(max_batch_size, save_engine):
        with trt.Builder(TRT_LOGGER) as builder, \
            builder.create_network(1 << (int)(trt.NetworkDefinitionCreationFlag.EXPLICIT_BATCH)) as network, \
            trt.OnnxParser(network, TRT_LOGGER) as parser:
            builder.max_workspace_size = 1 << 30
            builder.max_batch_size = max_batch_size

            builder.fp16_mode = False
            builder.int8_mode = False
            if not os.path.exists(onnx_file_path):
                quit('ONNX file {} not found'.format(onnx_file_path))
            print('Loading ONNX file from path {}'.format(onnx_file_path))
            with open(onnx_file_path, 'rb') as model:
                print('Beginning ONNX file parsing')
                if not parser.parse(model.read()):
                    for error in range(parser.num_errors):
                        print(parser.get_error(error))
            last_layer = network.get_layer(network.num_layers - 1)
            print('Completed parsing of ONNX file')
            print('Building an engine from file {}; this may take a while...'.format(onnx_file_path))

            engine = builder.build_cuda_engine(network)
            print('Completed creating Engine')

            if save_engine:
                with open(engine_file_path, "wb") as f:
                    f.write(engine.serialize())
                return engine
        if os.path.exists(engine_file_path):
            print('Reading engine from file {}'.format(engine_file_path))
            with open(engine_file_path, "rb") as f, trt.Runtime(TRT_LOGGER) as runtime:
                return runtime.deserialize_cuda_engine(f.read())
        else:
            return build_engine(max_batch_size, save_engine)

    def do_inference(context, bindings, inputs, outputs, stream, batch_size=1):
        [cuda.memcpy_htod_async(inp.device, inp.host, stream) for inp in inputs]
        context.execute_async(batch_size=batch_size, bindings=bindings, stream_handle=stream.handle)
        [cuda.memcpy_dtoh_async(out.host, out.device, stream) for out in outputs]
        stream.synchronize()
        return [out.host for out in outputs]

    def postprocess_the_outputs(h_outputs, shape_of_output):

```

```

h_outputs = h_outputs.reshape(*shape_of_output)
return h_outputs

def main():
    img = np.ones([1,3,224,224]) * 0.5
    img = img.astype(dtype=np.float32)

    fp16_mode = False
    int8_mode = False
    engine = get_engine(max_batch_size, onnx_model_path, trt_engine_path, save_engine=True)

    context = engine.create_execution_context()

    inputs, outputs, bindings, stream = allocate_buffers(engine)

    shape_of_output = (max_batch_size, 1000)
    inputs[0].host = img.reshape(-1)

    t1 = time.time()
    trt_outputs = do_inference(context, bindings=bindings, inputs=inputs, outputs=outputs, stream=stream)
    t2 = time.time()
    output_trt = postprocess_the_outputs(trt_outputs[0], shape_of_output)
    print('TensorRT OK!')
    # print(np.argmax(output_trt))

    model = torchvision.models.alexnet(pretrained=True).cuda()
    alexnet_model = model.eval()

    input_for_torch = torch.from_numpy(img).cuda()
    t3 = time.time()
    output_torch = alexnet_model(input_for_torch)
    t4 = time.time()
    output_torch = output_torch.cpu().data.numpy()
    print('Pytorch OK!')
    # print(np.argmax(output_torch))

    mae = np.mean(abs(output_trt - output_torch))
    print('Inference time with the TensorRT engine: {}'.format(t2-t1))
    print('Inference time with the Pytorch model: {}'.format(t4-t3))
    print('MAE = {}'.format(mae))

    print('All completed!')

if __name__ == '__main__':
    main()

```

- 运行结果

```

Reading engine from file alexnet.trt
[TensorRT] WARNING: Current optimization profile is: 0. Please ensure there are no enqueued operations pending in this context prior to switching profiles
[TensorRT] WARNING: Explicit batch network detected and batch size specified, use enqueue without batch size instead.
TensorRT OK!
Pytorch OK!
Inference time with the TensorRT engine: 0.004355669021606445
Inference time with the Pytorch model: 0.005367755889892578
MAE = 3.7649647310900036e-07
All completed!

```

4. 注意事项

- tensorrt只支持zero pad

<https://github.com/NVIDIA/TensorRT/issues/195>

- tensorrt 7.0.0.11不支持batchnorm层，可以从 [doc/pdf/TensorRT-Developer-Guide.pdf](#) (Appendix A.1. TensorRT Layers章节) 查看tensorrt支持的层